

CSC 374/407: Computer Systems II

Lecture 5

Joseph Phillips
De Paul University

2017 February 11

Copyright © 2011-2017 Joseph Phillips
All rights reserved

Reading

- ♦ Bryant & O'Hallaron “*Computer Systems, 3rd Ed.*”
 - Chapter 12.3-12.7: Concurrent Programming
- ♦ Hoover “*System Programming*”
 - (None)

Topics

- ◆ What threads are and why use them
- ◆ How to create threads
- ◆ Critical sections and unsafe thread programming
- ◆ Synchronization
- ◆ Application: Producer/Consumer

What threads are and why use them

Purpose: to have multiple “little processes”
parallelize action of one program

Threads are like processes, but they share:

- Code segment
- Heap segment
- Data/BSS segments
- Stack segment

They differ in that:

- Have different registers (including condition codes)
- Start in different regions of stack
- Have unique thread id

How to compile for threads

At the beginning of your program:

```
. . .  
#include <thisHeader.h>  
#include <pthread.h>  
#include <thatHeader.h>  
. . .
```

When you link:

```
unix> gcc . . . -lpthread . . .
```

Why are they called “p-threads?”

- **P** = **POSIX** = **P**ortable **O**perating **S**ys **I**nterface for **UniX**
- Standardizes “Unix” across Linux, BSD, Solaris, *etc.*

A (barely) threaded program

```
#include<stdlib.h>
#include<stdio.h>
#include<pthread.h>
```

Don't forget, now!

```
struct TwoInts { int int0;  int int1; };
```

```
void* fnc  (void* vPtr);
```

```
int main  ()
```

```
{ pthread_t      tId;
```

```
  struct TwoInts  twoInts;
```

```
  twoInts.int0 = 2;  twoInts.int1 = 3;
```

```
  pthread_create(&tId,NULL,fnc,(void*)&twoInts);
```

```
  int*          intPtr;
```

```
  pthread_join(tId, (void**)&intPtr);
```

```
  printf("The sum is %d\n",*intPtr);
```

```
  return(EXIT_SUCCESS);
```

```
}
```

A (barely) threaded program

```
#include<stdlib.h>
#include<stdio.h>
#include<pthread.h>
```

**pthread_create needs the
addr of a thread object in
which to build the thread**

```
struct TwoInts { int int0;  int int1; };
```

```
void* fnc  (void* vPtr);
```

```
int main  ()
```

```
{ pthread_t      tId;
```

```
  struct TwoInts  twoInts;
```

```
  twoInts.int0  = 2;  twoInts.int1  = 3;
```

```
  pthread_create(&tId, NULL, fnc, (void*)&twoInts);
```

```
  int*          intPtr;
```

```
  pthread_join(tId, (void**)&intPtr);
```

```
  printf("The sum is %d\n", *intPtr);
```

```
  return(EXIT_SUCCESS);
```

```
}
```

A (barely) threaded program

```
#include<stdlib.h>
#include<stdio.h>
#include<pthread.h>
```

You may create a variety of different threads. **NULL** means "an ordinary thread"

```
struct TwoInts { int int0;  int int1; };
```

```
void* fnc  (void*  vPtr);
```

```
int main  ()
```

```
{ pthread_t  tId;
```

```
  struct TwoInts  twoInts;
```

```
  twoInts.int0  =  2;  twoInts.int1  =  3;
```

```
  pthread_create(&tId,NULL,fnc,(void*)&twoInts);
```

```
  int*          intPtr;
```

```
  pthread_join(tId, (void**)&intPtr);
```

```
  printf("The sum is %d\n",*intPtr);
```

```
  return(EXIT_SUCCESS);
```

```
}
```


A (barely) threaded program

```
#include<stdlib.h>
#include<stdio.h>
#include<pthread.h>
```

The thread runs a different
fnc in the same program.
Give fnc name, not fnc call

```
struct TwoInts { int int0; int int1; };
```

```
void* fnc (void* vPtr);
```

```
int main ()
```

```
{ pthread_t tId;
```

```
struct TwoInts twoInts;
```

```
twoInts.int0 = 2; twoInts.int1 = 3;
```

```
pthread_create(&tId, NULL, fnc, (void*)&twoInts);
```

```
int* intPtr;
```

```
pthread_join(tId, (void**)&intPtr);
```

```
printf("The sum is %d\n", *intPtr);
```

```
return(EXIT_SUCCESS);
```

```
}
```

A (barely) threaded program

```
#include<stdlib.h>
#include<stdio.h>
#include<pthread.h>
```

**fnc() takes void* arg. void*
can be address of anything,
like Java's Object.**

```
struct TwoInts { int int0;  int int1; };
```

```
void* fnc  (void* vPtr);
```

```
int main  ()
```

```
{ pthread_t      tId;
```

```
  struct TwoInts  twoInts;
```

```
  twoInts.int0 = 2;  twoInts.int1 = 3;
```

```
  pthread_create(&tId, NULL, fnc, (void*)&twoInts);
```

```
  int*      intPtr;
```

```
  pthread_join(tId, (void**)&intPtr);
```

```
  printf("The sum is %d\n", *intPtr);
```

```
  return(EXIT_SUCCESS);
```

```
}
```

A (barely) threaded program

```
#include<stdlib.h>
#include<stdio.h>
#include<pthread.h>
```

pthread_join() waits for particular thread, like **wait_pid()**. No '&'.

```
struct TwoInts { int int0; int int1; };
```

```
void* fnc (void* vPtr);
```

```
int main ()
```

```
{ pthread_t tId;
```

```
struct TwoInts twoInts;
```

```
twoInts.int0 = 2; twoInts.int1 = 3;
```

```
pthread_create(&tId, NULL, fnc, (void*)&twoInts);
```

```
int* intPtr;
```

```
pthread_join(tId, (void**)&intPtr);
```

```
printf("The sum is %d\n", *intPtr);
```

```
return(EXIT_SUCCESS);
```

```
}
```

A (barely) threaded program

```
#include<stdlib.h>
#include<stdio.h>
#include<pthread.h>
```

This one is funky,
and needs further
explanation . . .

```
struct  TwoInts { int int0;  int int1; };
```

```
void* fnc  (void*  vPtr);
```

```
int main  ()
```

```
{ pthread_t      tId;
```

```
  struct TwoInts  twoInts;
```

```
  twoInts.int0 = 2;  twoInts.int1 = 3;
```

```
  pthread_create(&tId,NULL,fnc,(void*)&twoInts);
```

```
  int*          intPtr;
```

```
  pthread_join(tId, (void**)&intPtr);
```

```
  printf("The sum is %d\n",*intPtr);
```

```
  return(EXIT_SUCCESS);
```

```
}
```

A (barely) threaded program


```
/* By The Way, here is the function the thread  
will run */
```

```
void* fnc (void* vPtr)  
{  
    static int      sum;  
    struct TwoInts* twoIntsPtr;  
  
    twoIntsPtr = (struct TwoInts*)vPtr;  
    sum = twoIntsPtr->int0 + twoIntsPtr->int1;  
    return( (void*)&sum );  
}
```


*(void**) &Ptr?!? WTF?!?*

Question 1: Why pointers?

Sometimes the thread wants to **return a tiny object**:



Sometimes the thread wants to **return a huge object**:



?? How can we handle this difference in size ??

*(void**) &Ptr?!? WTF?!?*


Question 1: Why pointers?

A: By returning a **pointer** (silly)!

Q: *Hmm, how should we do it?*

*(void**) &Ptr?!? WTF?!?*

Question 2: Why &ptr?



// Attempt #1:

```
pthread_join(tId, intPtr)
```

// means . . .

```
pthread_join(tId, 0x3BFC)
```


// means . . .

“Make garbage at `0x3BFC` point to 5”

// means . . .

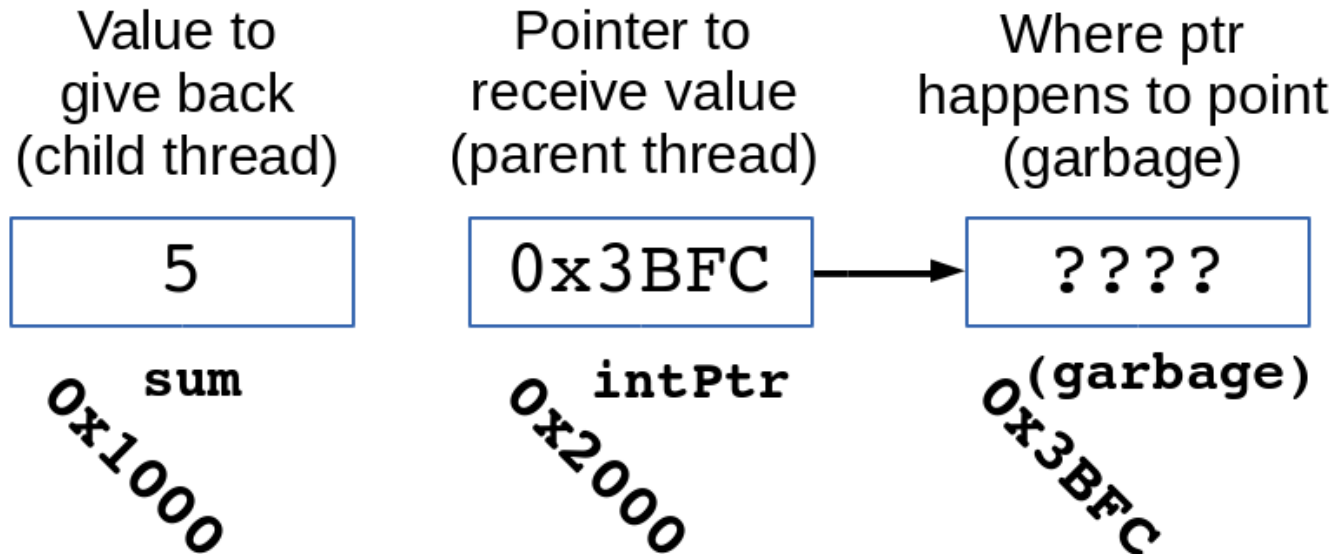
*(void**) &Ptr?!? WTF?!?*

Question 2: Why &ptr?



*(void**) &Ptr?!? WTF?!?*

Question 2: Why &ptr?



// Attempt #2:

```
pthread_join(tId, &intPtr)
```

// means . . .


```
pthread_join(tId, 0x2000)
```

// means . . .

"Make intPtr point to 5"

// means . . .

*(void**)&Ptr?!? WTF?!?*
Question 2: Why &ptr?



*(void**) &Ptr?!? WTF?!?*

Ques 3: Why (void)&ptr?**

Q: *So I see why you say **&ptr**, why **(void**)**?*

A: **(void*)** means *just an address* (by making it **void** you have forgotten object type) (like **Object** in Java)

(void)** means the address of a *pointer* that can hold *just an address*.

Detached threads

Called “detached” because no need to `pthread_join()` with parent thread.

```
pthread_t      tId;
pthread_attr_t tAttr;

pthread_attr_init(&tAttr);
pthread_attr_setdetachstate(&tAttr, PTHREAD_CREATE_DETACHED);
pthread_create(&tId, &tAttr, fncName, &args);
...
pthread_attr_destroy(&tAttr);
```

Good for servers:

- Parent thread makes new thread for each client
- Parent does not have to `pthread_join()`

How to create threads (in summary)

```
int pthread_create
(pthread_t* restrict      thread,
 /* Pointer to a pthread_t object to identify
  the child thread */
const pthread_attr_t* restrict attr,
 /* Pointer to optional object for properties
  of child.  You can just say NULL. */
void *(*fncName)(void*),
 /* Name of function to run:
  void* fncName(void* ptr) */
void *restrict           arg
 /* Ptr to object that is arg to fncName() */
)
```

Return value is:

- 0: success
- anything else: ERROR!

How to wait for threads (in summary)

```
int pthread_join
    (pthread_t          thread,
     /* Thread for which to wait */
     void**             valuePtr
     /* Value returned by pthread_exit()
        pointed to by valuePtr (may be NULL)
        */
    )
```

Compare threads with fork:

```
/* I use fork() to create a process */
```

```
#include <stdlib.h>
```

```
#include <stdio.h>
```

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
int main()
```

```
{
```

```
    int  childStatus;
```

```
    pid_t childId  = fork();
```

```
    char* argsText = /* whatever */
```

```
    if (childId == 0)
```

```
        execlp("doThisFile",...,argsText,.. );
```

```
    waitpid(childId,&childStatus,0);
```

```
    return(0);
```

```
}
```

```
/* I use pthread_create() */
```

```
#include <stdlib.h>
```

```
#include <stdio.h>
```

```
#include <pthread.h>
```

```
int main()
```

```
{
```

```
    pthread_t      childId;
```

```
    SendObject      arg = /*whatever*/
```

```
    ReceiveObject*  childStatusPtr;
```

```
    pthread_create(&childId,NULL,  
                  doThisFnc,&arg  
                  );
```

```
    pthread_join(childId,&childStatusPtr);
```

```
    return(0);
```

```
}
```


Example program:

```
/* Compile with:
 * gcc -lpthread thread_ex1.c -o thread_ex1
 */
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <pthread.h>
const int N = 2;

/* thread routine */
void *thread_routine(void *vargp)
{
    int id = *(int*)vargp;
    char* ptr;
    printf("Hello from child thread %d\n", id);
    switch (id)
    {
        case 0: ptr = strdup("Hello "); break;
        case 1: ptr = strdup("there!"); break;
    }
    return(ptr);
}
```

```
int main()
{
    int i;
    char* msgPtr;
    pthread_t tid[N];

    for (i = 0; i < N; i++)
        pthread_create
            (&tid[i], NULL,
             thread_routine, (void *)&i);

    for (i = 0; i < N; i++)
    {
        pthread_join(tid[i], (void**)&msgPtr);
        puts(msgPtr);
        free(msgPtr);
    }

    return(0);
}
```

Time for you!

Can you write a threaded program that can:

1. **Prove** that all threads use the same **stack** (did we just do that?)
2. **Prove** that all threads use the same **global var data space** (for global vars and static vars inside functions)
3. **Prove** that all threads use the same **heap**.

What's wrong with this?

```
/* Compile with:
 * gcc -lpthread badcnt.c -o badcnt
 */
#include <stdlib.h>
#include <stdio.h>
#include <pthread.h>

unsigned int cnt = 0; /* shared */
unsigned int NUM_ITERS;

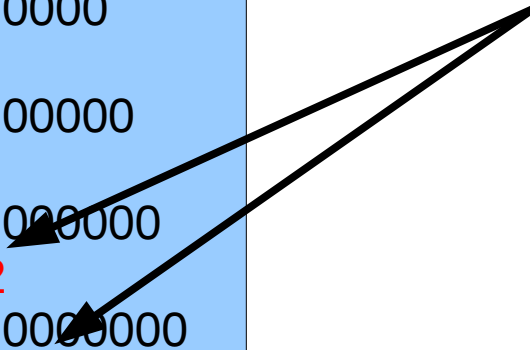
/* thread routine */
void *count(void *arg)
{
    int i;
    for (i=0; i<NUM_ITERS; i++)
        cnt++;
    return NULL;
}
```

```
int main(int argc, char* argv[])
{
    pthread_t tid1, tid2;
    if ( (argc >= 2) && isdigit(*argv[1]) )
        NUM_ITERS = atoi(argv[1]);
    else
    {
        const int LINE_SIZE = 16;
        char line[LINE_SIZE];
        do {
            printf("How many iterations? ");
            fgets(line, LINE_SIZE, stdin);
        }
        while ( !isdigit(line[0]) );
        NUM_ITERS = atoi(line);
    }
    pthread_create(&tid1, NULL, count, NULL);
    pthread_create(&tid2, NULL, count, NULL);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("Should be %d is %d\n", NUM_ITERS*2, cnt);
    return(0);
}
```

Output:

```
[jphillips@localhost]$ badcnt 10  
Should be 20 is 20  
[jphillips@localhost]$ badcnt 100  
Should be 200 is 200  
[jphillips@localhost]$ badcnt 1000  
Should be 2000 is 2000  
[jphillips@localhost]$ badcnt 10000  
Should be 20000 is 20000  
[jphillips@localhost]$ badcnt 100000  
Should be 200000 is 200000  
[jphillips@localhost]$ badcnt 1000000  
Should be 2000000 is 1830352  
[jphillips@localhost]$ badcnt 10000000  
Should be 20000000 is 15214384
```

Last two
values are
WRONG!



Critical sections

(gdb) disass count

```
0x400688: push  %rbp
0x400689: mov   %rsp,%rbp
0x40068c: mov   %rdi,0xffffffffffe8(%rbp)
0x400690: movl  $0x0,0xffffffffffc(%rbp)
0x400697: jmp   0x4006ac <count+36>
0x400699: mov   2098629(%rip),%eax # 0x600c64 <cnt>
0x40069f: add   $0x1,%eax
0x4006a2: mov   %eax,2098620(%rip) # 0x600c64 <cnt>
0x4006a8: addl  $0x1,0xffffffffffc(%rbp)
0x4006ac: mov   0xffffffffffc(%rbp),%edx
0x4006af: mov   2098611(%rip),%eax #0x600c68 <NUM_ITERS>
0x4006b5: cmp   %eax,%edx
0x4006b7: jb    0x400699 <count+17>
0x4006b9: mov   $0x0,%eax
0x4006be: leaveq
0x4006bf: retq
```



These
three
must
be done
by one
thread at
a time!

Critical sections (2)

Imagine some scenarios where they are **not** done atomically

Other ways threads can step on each other's toes

1. Working with the same global variable
 - (Just covered)
2. Working with the same static var within a fnc.
 - Can *you* think of an example?
3. A function returning a pointer to a static var.
 - Can *you* think of an example?
4. Calling a thread-unsafe function.
 - Can *you* think of an example?

Functions: safe and unsafe

Reentrant functions call no shared variables

- Always thread safe

Most standard C Library functions are thread safe

- e.g. printf(), malloc()

Unsafe to safe

- | | |
|---|---|
| – asctime() (<i>unsafe!</i>) | asctime_r() (<i>safe!</i>) |
| – ctime() (<i>unsafe!</i>) | ctime_r() (<i>safe!</i>) |
| – gethostbyaddr() (<i>unsafe!</i>) | gethostbyaddr_r() (<i>safe!</i>) |
| – inet_ntoa(<i>unsafe!</i>) | <i>NO SAFE VERSION!</i> |
| – localtime() (<i>unsafe!</i>) | localtime_r() (<i>safe!</i>) |
| – rand() (<i>unsafe!</i>) | rand_r() (<i>safe!</i>) |

The semaphore solution

Classic solution: Dijkstra's P and V operations on semaphores.

- **semaphore:** non-negative integer synchronization variable.
 - `sem_wait(s): [while(s==0) wait(); s--;]`
 - Originally named P(), Dutch for "Proberen" (test)
 - `sem_post(s): [s++;]`
 - Originally named V(), Dutch for "Verhogen" (increment)
- OS guarantees that operations between brackets `[]` are executed indivisibly.
 - Only one P or V operation at a time can modify s.
 - When while loop in P terminates, only that P can decrement s.

Semaphore invariant: $(s \geq 0)$

(1) POSIX Semaphores

What to include:

- `#include <semaphore.h>`

Types and functions:

- `sem_t semaphore;`
- `sem_init(sem_t* semPtr, int flag, int value)`
 - Initialize pointed-to semaphore, with `value`, if `flag == 1` then semaphore can be forked
- `sem_destroy(sem_t* semPtr)`
 - Destroy pointed-to semaphore. If it's negative then block.

POSIX Semaphores, cont'd

- `sem_wait(sem_t* semPtr)`
 - Decrement pointed-to semaphore. If it's negative then block.
- `sem_post(sem_t* semPtr)`
 - Increment pointed-to semaphore. Wake one blocked process if any.
- `sem_getvalue(sem_t* semPtr, int* valuePtr)`
 - Get value of pointed to semaphore.

POSIX semaphore solution

```
...
sem_t      sem;
...

/* thread routine */
void *count(void *arg)
{
    int i;
    for (i=0; i<NUM_ITERS; i++)
    {
        sem_wait(&sem);
        cnt++;
        sem_post(&sem);
    }
    return NULL;
}
```

```
/* in main */
...
/* initialize sem to 1 */
if (sem_init(&sem, 0, 1) < 0)
{
    fputs("sem_init error\n",stderr);
    exit(EXIT_FAILURE);
}
...
if (sem_destroy(&sem) < 0)
{
    fputs("sem_destroy error\n",stderr);
    exit(EXIT_FAILURE);
}
```

(2) *pthread_mutex* solution

- `pthread_mutex_t lock`
- `pthread_mutex_init(addressOfLock, NULL)`
 - Makes a lock and initializes it to default values.
- `pthread_mutex_destroy(addressOfLock)`
 - Makes a lock and initializes it to default values.
- `pthread_mutex_lock(address of lock)`
 - Blocks thread until lock obtained, then obtains lock and blocks other threads until lock released.
- `pthread_mutex_unlock(address of lock)`
 - Releases lock allowing other threads to obtain it.

pthread_mutex solution

```
pthread_mutex_t cntLock;

/* thread routine */
void *count(void *arg) {
    for (int i=0; i<NUM_ITERS; i++) {
        pthread_mutex_lock(&cntLock);
        cnt++;
        pthread_mutex_unlock(&cntLock);
    }
    return NULL;
}

/* in main() */
...
pthread_mutex_init(&cntLock,NULL);
...
pthread_mutex_destroy(&cntLock);
```

```
[jphillips@localhost]$ bettercnt 10
Should be 20 is 20
[jphillips@localhost]$ bettercnt 100
Should be 200 is 200
[jphillips@localhost]$ bettercnt 1000
Should be 2000 is 2000
[jphillips@localhost]$ bettercnt 10000
Should be 20000 is 20000
[jphillips@localhost]$ bettercnt 100000
Should be 200000 is 200000
[jphillips@localhost]$ bettercnt 1000000
Should be 2000000 is 2000000
[jphillips@localhost]$ bettercnt 10000000
Should be 20000000 is 20000000
[jphillips@localhost]$ bettercnt 100000000
Should be 200000000 is 200000000
```

Producer-Consumer

One or more threads produce something

- Place in buffer

One or more threads consume something

- Retrieve from buffer

Critical section

- Access of pointers/indices for buffer

Problem!

- A **producer** may gain access to a **full buffer**
- A **consumer** may gain access to an **empty buffer**

Oh no! An example in **C++**!

- A buffer is an **object**.
- Make the **object** thread-safe.
- Lessons **carry over to** Java, C#, etc.

Unsafe_Buffer.h

```
class Buffer
{
    enum { SIZE = 16 };

    int array_[SIZE];
    int inIndex_;
    int outIndex_;
    int numItems_;

public :

    Buffer ()
    {
        inIndex_ = outIndex_ = numItems_ = 0;
    }

    ~Buffer ()
    {
    }

    int getNumItems () const
    { return(numItems_); }

    void putIn (int i)
    {
        while (getNumItems() >= SIZE)
        {
            printf("Full! Waiting!\n");
            usleep(10);
        }

        array_[inIndex_] = i;
        countArray[array_[inIndex_]]++;
        usleep(10 + rand() % 10);
        inIndex_++;
        numItems_++;
        if (inIndex_ >= SIZE)
            inIndex_ = 0;
    }
}
```


Unsafe_Buffer.h

```
int    pullOut ()
{
    while  (getNumItems() <= 0)
    {
        printf("Empty!  Waiting!\n");
        usleep(10);
    }

    countArray[array_[outIndex_]]--;
    int toReturn      = array_[outIndex_];
    usleep(10 + rand() % 10);

    outIndex_++;
    numItems_--;
    if  (outIndex_ >= SIZE)
        outIndex_ = 0;

    return(toReturn);
}
};
```

producerConsumer.cpp

```
#include          <stdlib.h>
#include          <stdio.h>
#include          <unistd.h>
#include          <pthread.h>
int*              countArray;
#include          "Unsafe_Buffer.h"

const int         NUM_INTEGERS_TO_BUFFER  = 0x1000;

void*             stuffIntegersIn (void*   vPtr)
{
    for (int i = 0;  i < NUM_INTEGERS_TO_BUFFER;  i++)
        ((Buffer*)vPtr)->putIn(i);

    return(NULL);
}
```

producerConsumer.cpp

```
void*    pullIntegersOut (void*   vPtr)
{
    for  (int i = 0;   i < NUM_INTEGERS_TO_BUFFER;   i++)
    {
        int j = ((Buffer*)vPtr)->pullOut();

        printf("Trial %d got %d.\n",i,j);
        fflush(stdout);
    }

    return(NULL);
}
```

producerConsumer.cpp

```
int      main      ()
{
    pthread_t      producer0;
    pthread_t      producer1;
    pthread_t      consumer0;
    pthread_t      consumer1;
    Buffer          buffer;
    countArray     =
        (int*)calloc(NUM_INTEGERS_TO_BUFFER, sizeof(int));

    pthread_create(&producer0, NULL, stuffIntegersIn, &buffer);
    pthread_create(&producer1, NULL, stuffIntegersIn, &buffer);
    pthread_create(&consumer0, NULL, pullIntegersOut, &buffer);
    pthread_create(&consumer1, NULL, pullIntegersOut, &buffer);
}
```

producerConsumer.cpp

```
pthread_join(producer1,NULL);
pthread_join(producer0,NULL);
pthread_join(consumer1,NULL);
pthread_join(consumer0,NULL);

for (int i = 0; i < NUM_INTEGERS_TO_BUFFER; i++)
    if (countArray[i] < 0)
        printf("%d was gotten too many times!\n",i);
    else
        if (countArray[i] > 0)
            printf("%d was put too many times!\n",i);

return(EXIT_SUCCESS);
}
```

Your turn!

(1) Run it as-is. *Any problems?*

Your turn!

(2) Make it thread-safe.

```
void myMethod ( )  
{  
    pthread_mutex_lock(&lock);  
    doCriticalSection();  
    pthread_mutex_unlock(&lock);  
}
```

NOTE:

- 1) Each object gets its **own** lock.
- 2) All threads accessing the same object use **same** lock

Any problems still?

Solution: *pthread_cond*!

- `pthread_cond_t cond`
- `pthread_cond_init(addressOfCondition, NULL)`
 - Makes a condition and initializes it to default values.
- `pthread_cond_destroy(addressOfCondition)`
 - Destroys condition.
- `pthread_cond_wait(addrOfCondition, addrOfLock)`
 - Blocks thread until lock released and condition signaled. Then obtains lock again.
- `pthread_cond_signal(addressOfCondition)`
 - Signal *one waiting thread* that condition is met.
- `pthread_cond_broadcast(addressOfCondition)`
 - Signal *all waiting threads* that condition is met.

The proper solution

```
void myMethod ( )
{
    pthread_mutex_lock(&lock);
    while ( !this->isReady() )
        pthread_cond_wait(&wCond,&lock);
    doCriticalSection();
    pthread_cond_signal(&sCond);
    pthread_mutex_unlock(&lock);
}
```

NOTE:

- (1) **isReady()** is a method of the class. It determines if this object is ready for thread to do **myMethod()**. (E.g. A buffer's **putIn()** waits for **!this->isFull()**)
- (2) Conditions **wCond** and **sCond** may or may not be same.

We have mutual exclusion, are guaranteed to be safe?

Sorry but No!

- Race conditions
 - When result is dependent on order of processes or threads
- Deadlock
 - Thread/process A has lock 1 and is waiting for lock 2
 - Thread/process B has lock 2 and is waiting for lock 1

We don't have time now, but also check out

- ▶ **shmget () , shmat ()** : A way of having multiple processes (e.g. parent and child, or children of same parent) share memory, like how threads can.

Next time: *Memory!*