

CSC 374/407: Computer Systems II

Lecture 8

Joseph Phillips
De Paul University

2016 September 6

Copyright © 2012-2016 Joseph Phillips
All rights reserved

Reading

- ♦ Bryant & O'Hallaron “*Computer Systems, 3rd Ed.*”
 - Chapter 10: System Level I/O
 - Chapter 11: Networking Programming
- ♦ Hoover “*System Programming*”
 - Chapter 5: Input/Output

Topics

Unix Filesystem Design: A Process' Prospective

Unix Filesystem Design: A Systemwide Prospective

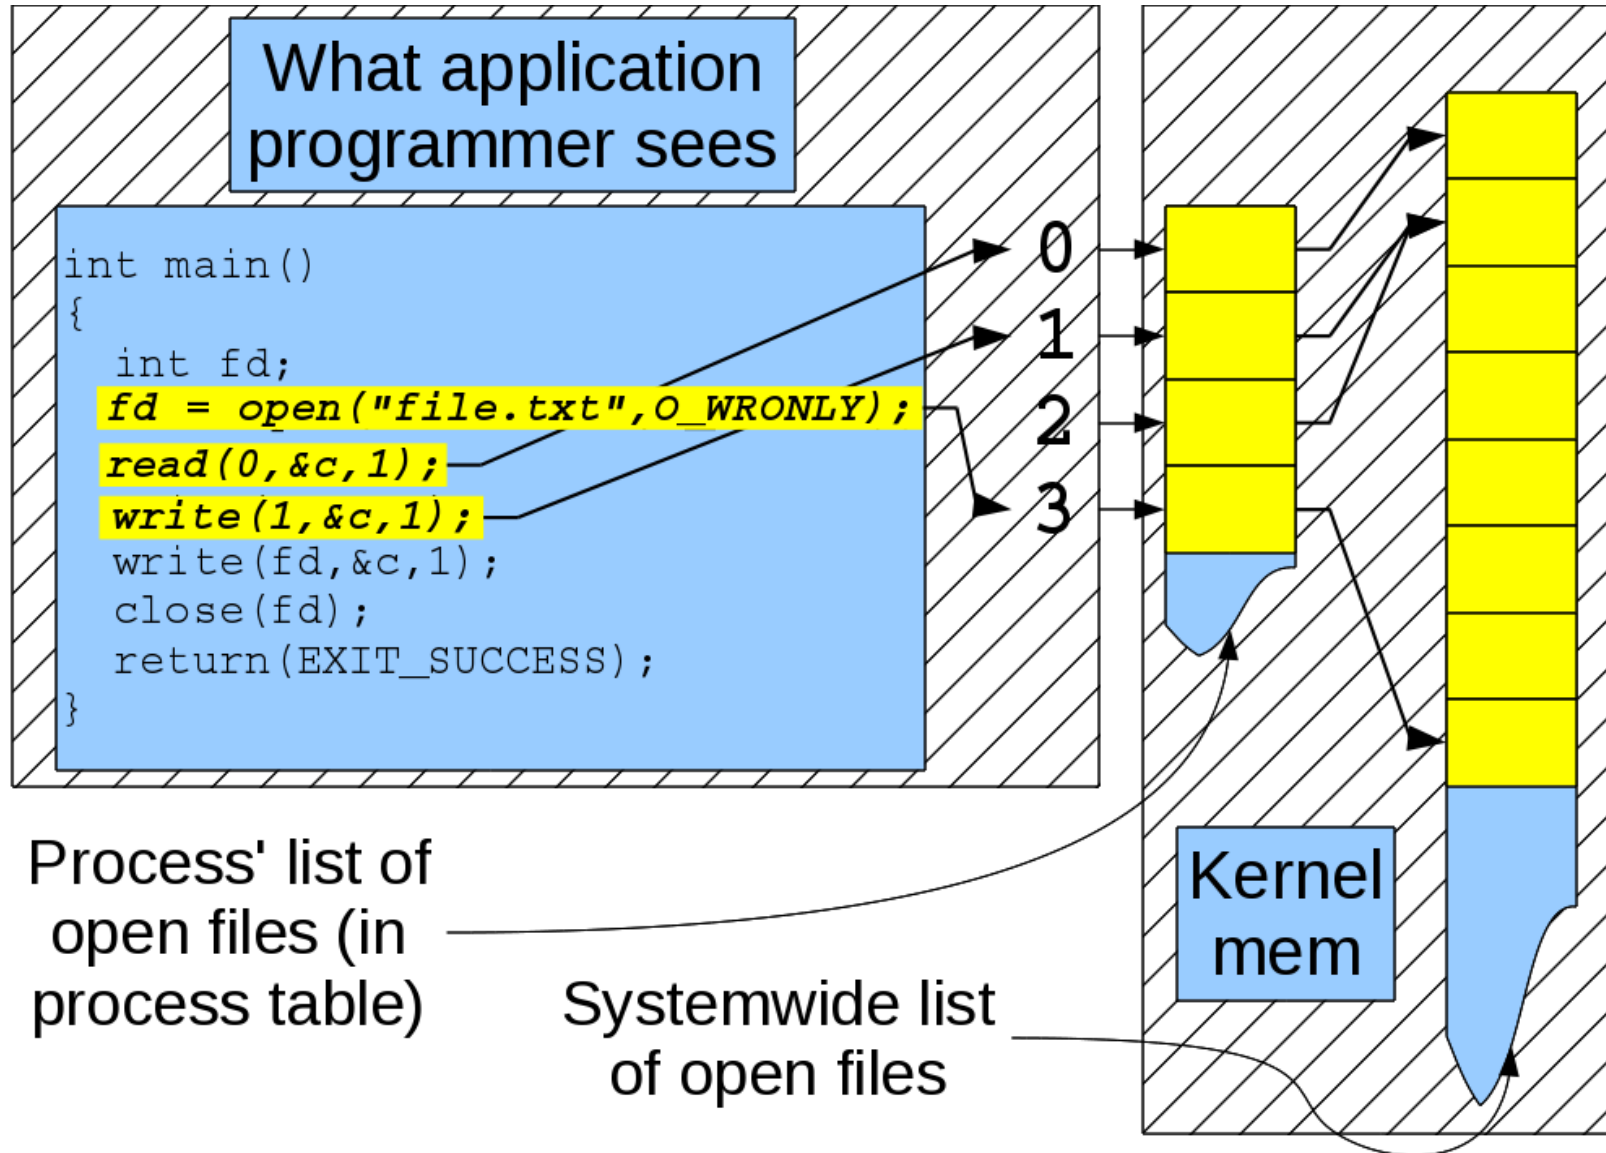
Low-level C Input-Output

Socket communication and the client/server model

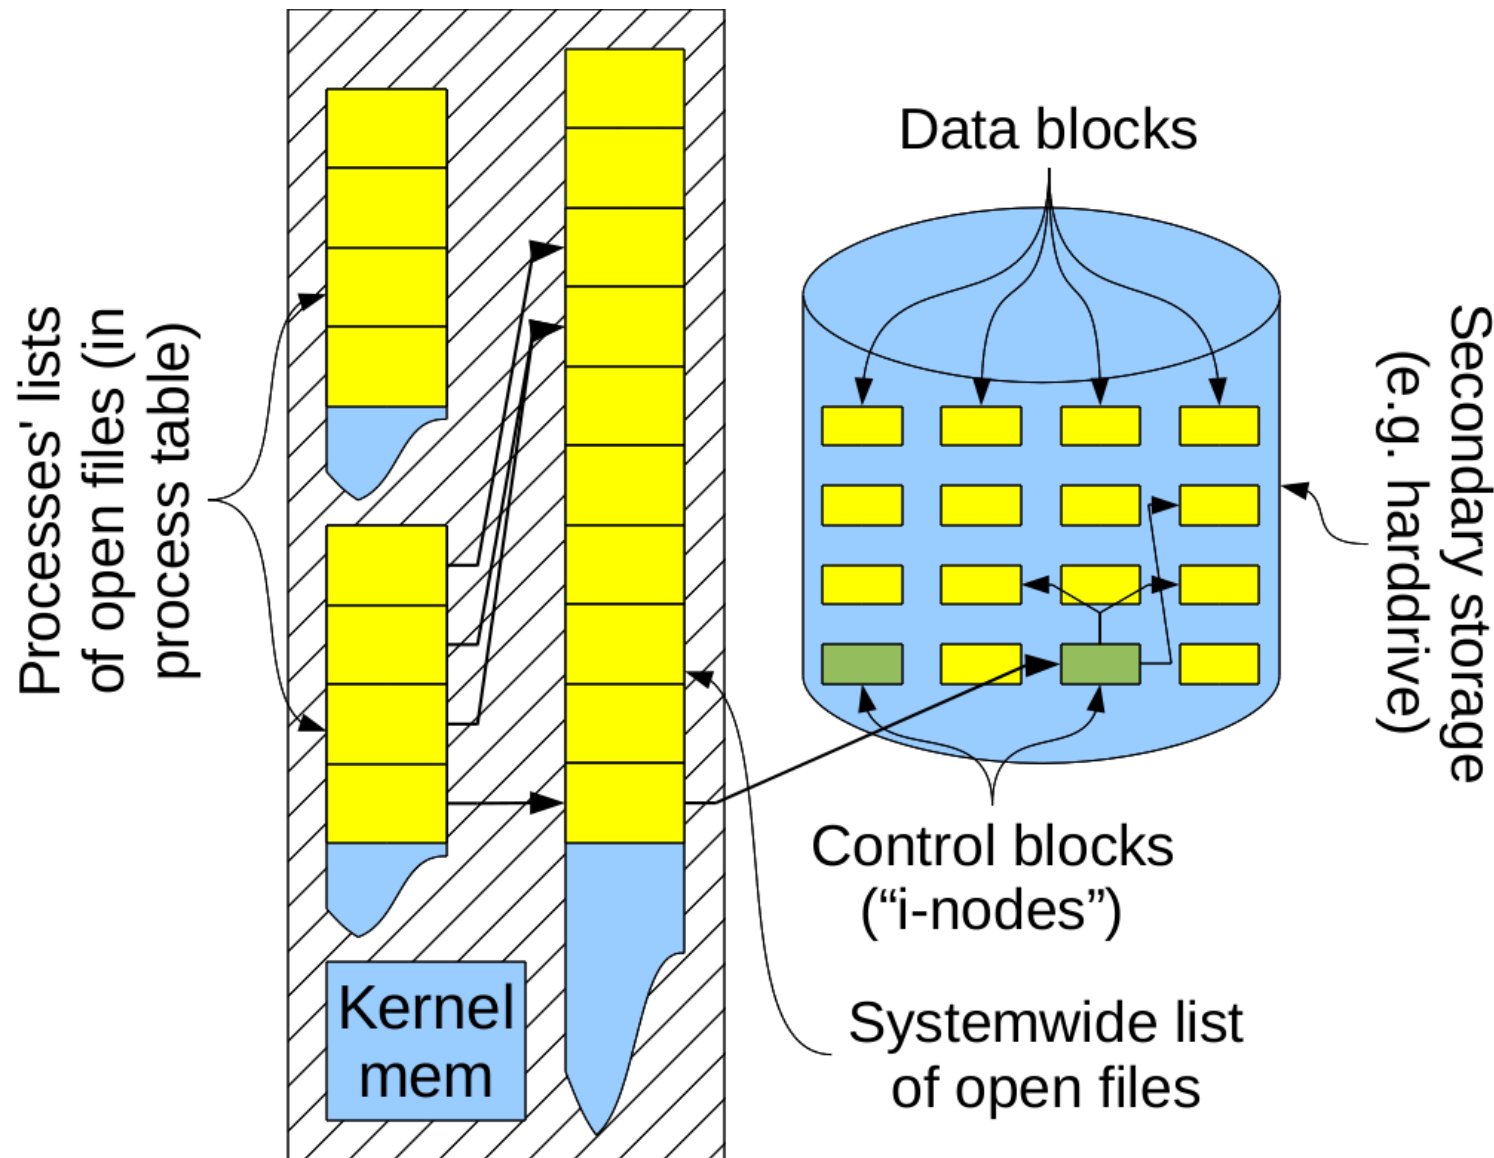
Server-side socket programming

Client-side socket programming

Unix Filesystem Design: A Process' Prospective



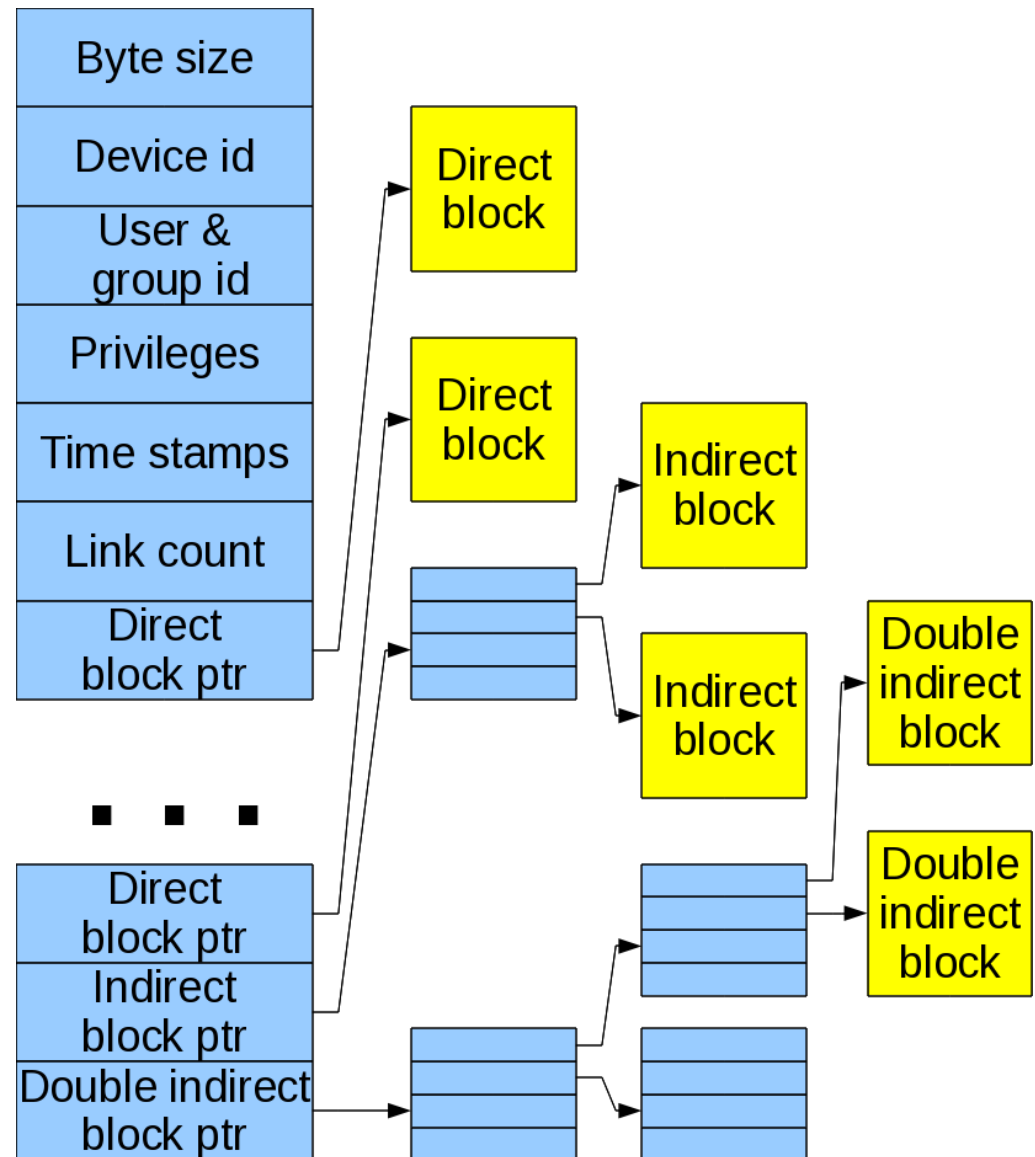
Unix Filesystem Design: A Systemwide Prospective



What's an “I-Node”?

Tells a files:

- Size in bytes
- Access times (last read, last written, last its status was modified)
- User and group ID
- Device ID
- Access privileges
- Link count (num different names/directories)
- Pointers



Low level C Input-Output

File descriptors are indices into process' file table

- 0: Standard input (*stdin*)
- 1: Standard output (*stdout*)
- 2: Standard error (*stderr*)

Useful commands include:

```
int open(const char* path, int how,  
int permission)
```

```
int close(int fd)
```

```
int read(int fd, char* bufferPtr,  
size_t bufferSize)
```

```
int write(int fd, char* bufferPtr,  
size_t numBytes)
```

```
int dup(int fd);
```

```
int pipe(int** );
```

open ()

int open(const char path, int how,
int permission)*

- Returns file descriptor (index into process' file array)
- File path given by path.

open ()

```
int open(const char* path, int how,  
int permission)
```

- Integer **how** is bitwise or-ing of one of:
 - **O_RDONLY**: Open for reading only.
 - **O_WRONLY**: Open for writing only.
 - **O_RDWR**: Open for reading and writing.
- And perhaps one or more of:
 - **O_CREAT**: Create file if doesn't already exist
 - **O_TRUNC**: If exist truncate its length to 0 (even if not open for writing)
 - **O_EXCL**: If **O_CREAT** is also set fail if file exists.
 - **O_APPEND**: Write to end of file.

open() , cont'd

```
int open(const char* path, int flags,  
int permission)
```

- In “Classic” Unix there are 3 types of permissions:
 - “**R**”ead: permission to read (i.e. load a file into memory)
 - “**W**”rite: permission to over-write a file
 - e“**X**”ecute: permission to run an executable
- Each of these permissions is either “On” (1) or “Off” (0)
- This suggests **storing all three** as an **octal digit**:
 - $\text{Read} * 2^2 + \text{Write} * 2^1 + \text{Execute} * 2^0$
- Examples:
 - Read and write, but no execute: $1*4 + 1*2 + 0*1 = 6$
 - Read and execute, but no write: $1*4 + 0*2 + 1*1 = 5$

open () , cont 'd

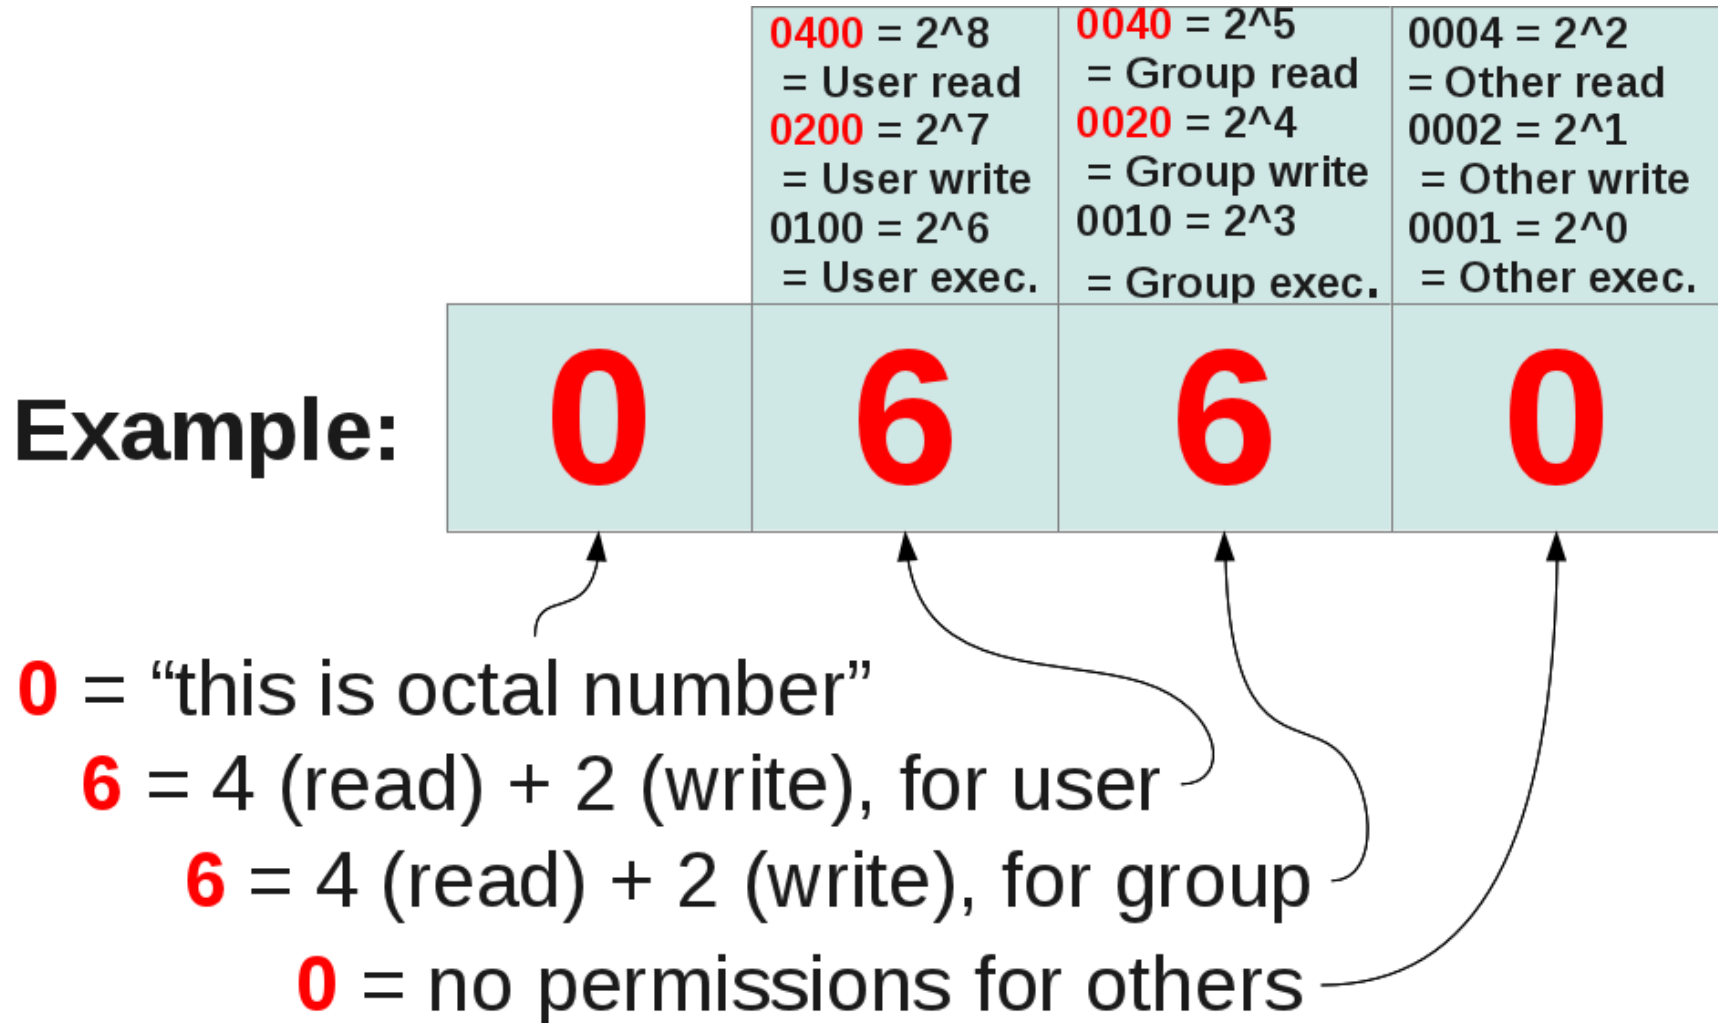
```
int open(const char* path, int flags,  
int permission)
```

Further, three types of folks for whom must specify permission:

- “User” (i.e. the Owner) of the file
- “Group”, audience of file
 - May just be User, or may be group to which user belongs, like CSC374Students
- “Other”, anyone else on OS.

Put it all together as a 3-digit Octal number
(Well 4-digit, leading “0” means “Octal”)

open() , cont'd



YOUR TURN

(1) What is the **octal code** for:
user: read, write;
group: read-only;
other: read-only

(2) What does octal code **0750** allow the
user, group and everyone else
to do?

close()

```
int close (int fd)
```

- Closes `fd`.
- Returns 0 on success or `-1` otherwise.
- Does not flush file.

Your turn!

You are going to open a process' *first file* after standard input (0), standard output (1) and standard error (2). The file descriptor is an index in a table. What is its integer value?

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/stat.h>
#include <fcntl.h>

int main ()
{
    int fd = open("bubu.txt", O_WRONLY | O_CREAT | O_APPEND, 0660);
    printf("fd = %d\n", fd);
    close(fd);
    return(EXIT_SUCCESS);
}
```

`write()`

```
int write(int fd, char* bufferPtr,  
size_t numBytes)
```

- Writes `numBytes` pointed to by `bufferPtr` to `fd`.
- Returns number of bytes written, or `-1` on error.

write() example

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/stat.h>
#include <fcntl.h>

int main ()
{
    int fd = open("bubu.txt", O_WRONLY | O_CREAT | O_APPEND, 0660);
    write(fd, "Bubu!\n", 6);
    close(fd);
    system("ls -l ./bubu.txt");
    return(EXIT_SUCCESS);
}
```

read()

```
int read(int fd, char* bufferPtr,  
size_t bufferSize)
```

- Reads up to **bufferSize** bytes from **fd** and puts them into **bufferPtr**.
- Returns number of bytes read from file, either
 - 0 (*“No more left!”*),
 - **bufferSize** (*“Here's a whole buffer full!”*),
 - somewhere inbetween (*“Here's all that's left”*), or,
 - -1 (*“Error!”*)

read() example

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/stat.h>
#include <fcntl.h>

#define BUFFER_SIZE 256

int main ()
{
    char buffer[BUFFER_SIZE];
    int  fd = open("bubu.txt", O_RDONLY, 0660);

    read(fd, buffer, BUFFER_SIZE);
    printf("%s", buffer);
    close(fd);
    return(EXIT_SUCCESS);
}
```

Your turn!

Write your own simple version of the Unix **littleCopy** file copying command. I'll get you started:

```
#include <stdlib.h>
#include <stdio.h>
#include <ctype.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <string.h>

#define          BUFFER_SIZE          256

/* Continued on next slide */
```

Your turn!

```
/* From previous slide */

int main (int argc, const char* argv[])
{
    const char* fromFileCPtr;
    const char* toFileCPtr;

    if (argc < 3)
    {
        fprintf(stderr,
                "Usage: littleCopy <fromFile> <toFile>\n"
                );
        return(EXIT_FAILURE);
    }

    fromFileCPtr = argv[1];
    toFileCPtr   = argv[2];

    /* YOUR CODE HERE */
    return(EXIT_SUCCESS);
}
```

Your turn, again!

```
/* Write a program that counts the number of occurrences
   of a character given on the command line. */
int main (int argc, const char* argv[])
{
    const char  charToCount;
    const char* fileCPtr;

    if  (argc < 3)
    {
        fprintf(stderr,
                  "Usage: charCount <char> <file>\n"
                );
        return(EXIT_FAILURE);
    }

    charToCount = *argv[1];
    fileCPtr    = argv[2];

    /* YOUR CODE HERE */
    return(EXIT_SUCCESS);
}
```

And your turn, yet again!

Revise the previous program to
count the number of lines in a file.

What happens if mother and child write to same file?

```
#include <stdlib.h>
#include <stdio.h>
#include <ctype.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <string.h>

#define FILENAME "bubu.txt"

int main ()
{
    const char* wordsPtr;
    int i;
    int numBytes;
    int pid;
    int fd =
    open(FILENAME,
        O_WRONLY|O_CREAT|O_TRUNC,
        0660);

    if (fd < 0)
    {
        fprintf(stderr,
            "Sorry, I can't make "
            "the output file %s\n",
            FILENAME);
        return(EXIT_FAILURE);
    }

    pid = fork();

    if (pid < 0)
    {
        fprintf(stderr,
            "Too many processes ace!\n"
            );
        return(EXIT_FAILURE);
    }
```


What happens if mother and child write to same file?

```
else
if (pid == 0)
    wordsPtr =
        "Baby says \"Gaga Gugu!\"\n";
else
    wordsPtr =
        "Mama says \"Poor baby!\"\n";


numBytes = strlen(wordsPtr);

for (i = 0; i < 4; i++)
{
    sleep(1);
    write(fd, wordsPtr,
          numBytes);
    printf(wordsPtr);
}

if (pid > 0)
{
    sleep(2);
    close(fd);
}

return(EXIT_SUCCESS);
}
```

What's going on?



Hey! Maybe we can use this for interprocess communication!

```
#include <unistd.h>
. . .
const int PIPE_READ = 0;
const int PIPE_WRITE = 1;
int      myPipe[2];

if (pipe(myPipe) == 0)
{
    char myArray[6];
    write(myPipe[PIPE_WRITE], "Hello!", 6);
    read (myPipe[PIPE_READ ], myArray, 6);
}
```



myPipe: An OS-owned buffer

dup()

```
#include <stdlib.h>
#include <stdio.h>
#include <sys/stat.h>
#include <fcntl.h>
```

```
#define FILENAME    "bubu.txt"
```

```
int main ()
{
    in fd= open(FILENAME,
                O_WRONLY|O_CREAT|O_TRUNC,
                0660);

    close(1); // Close stdout
    dup(fd);  // Redirect stdout to FILENAME
    printf("I wonder where this will show up?\n");
    close(fd); // Be polite!
    return(EXIT_SUCCESS);
}
```

stdin	0
stdout	1
stderr	2
bubu.txt	3

stdin	0
stdout	1
stderr	2
bubu.txt	3

stdin	0
bubu.txt	1
stderr	2
bubu.txt	3

dup () copies the entry of the given file descriptor to the first free one.

1

2

3

1

IPC with pipes

```
#include <stdlib.h>
#include <stdio.h>
#include <ctype.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <string.h>

int      main      ()
{
    int    parentToChild[2];
    int    childToParent[2];

    if    ((pipe(parentToChild) < 0)
           || (pipe(childToParent) < 0))
    {
        fprintf(stderr,
                "Can't make pipes\n");
        return(EXIT_FAILURE);
    }

    int    pid      = fork();

    if    (pid < 0)
    {
        fprintf(stderr, "Too many
processes Ace!\n");
        return(EXIT_FAILURE);
    }
    else
    if    (pid == 0)
    {
        //  Baby's case
        close(0);    // Close "stdin"
        dup(parentToChild[0]);
        close(1);    // Close "stdout"
        dup(childToParent[1]);

        // . . . continued
    }
}
```

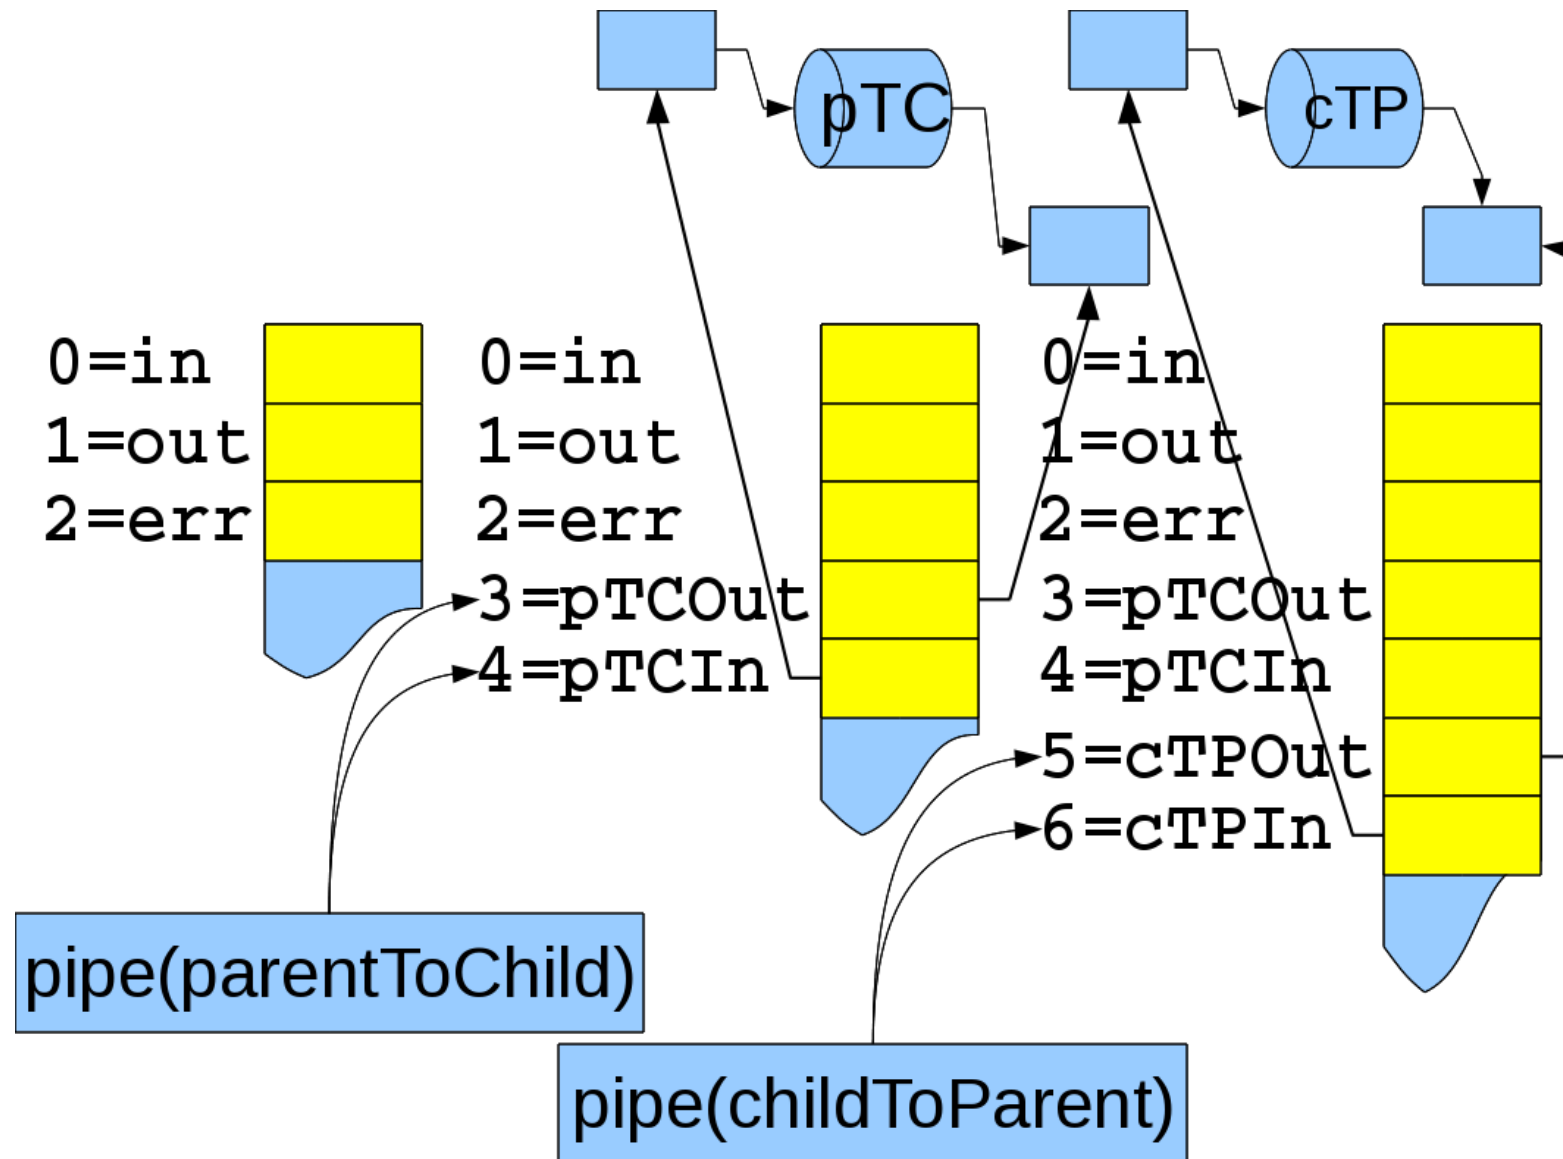
IPC with pipes, cont'd

```
// Baby's case, continued
while (1)
{
    char buffer[10];
    int i,numRead;
    numRead =
        read(0,buffer,10);


    for (i=0; i<numRead; i++)
        buffer[i] =
            toupper(buffer[i]);

    write(1,buffer,numRead);
}
else
{
    // Mama's case
    while (1)
    {
        char buffer[10];
        fgets(buffer,10,stdin);
        write
            (parentToChild[1],
             buffer,
             10);
        read
            (childToParent[0],
             buffer,
             10);
        printf(buffer);
    }
    return(EXIT_SUCCESS);
}
```


dup() and pipe(), 1




dup() and pipe(), 2



dup () and pipe (), 3



dup () and pipe (), 4



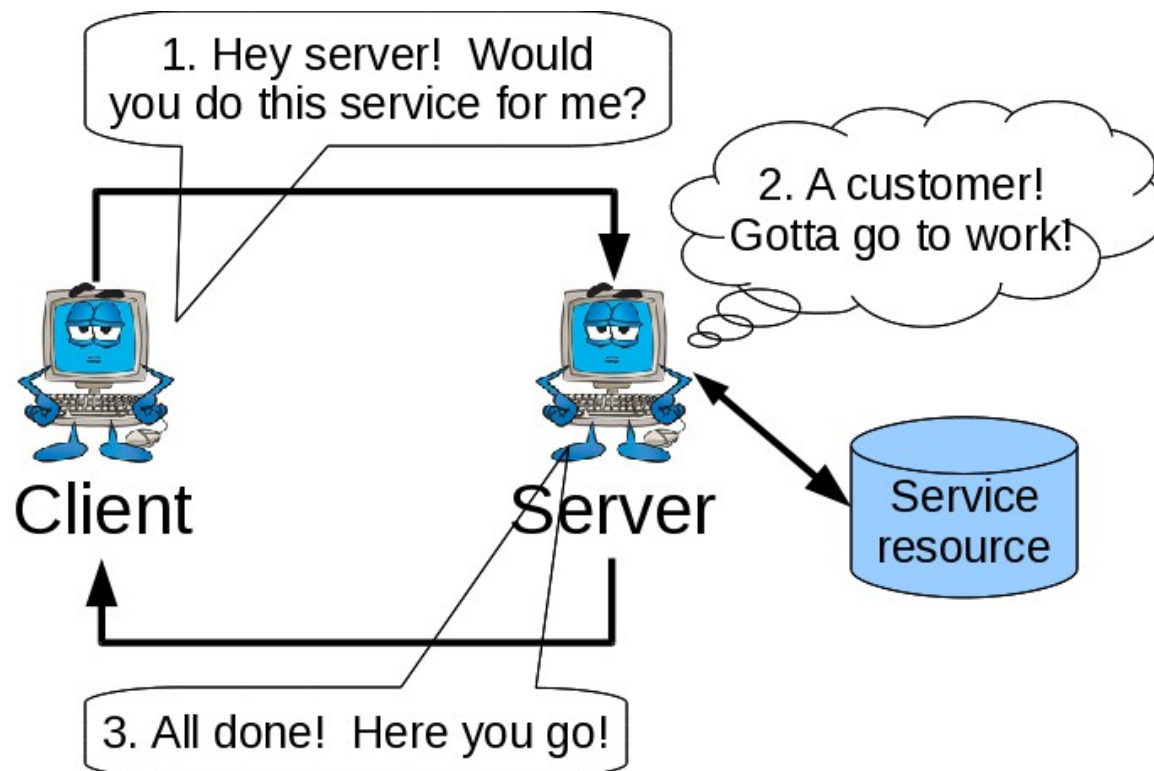
Sockets

- ◆ Hey! Moving bytes to/from a file descriptor is such a ***grand idea*** that we can use it to move bytes to/from ***another process***
- ◆ Further, the process could be ***here*** (on the same machine)
- ◆
..... or way ***over here***
(another machine across the Internet)
- ◆ We call this ***sockets!***

The client-server computing model

1. Client asks server for a service
2. Server does service
3. Server returns result to client

Examples: **ssh**, **sftp**, **http**, *etc.*



Processes talking to each other on different computers

Identify service by *IP address* and *port*

IP address: *Which computer?*

Humans like strings: “*www.depaul.edu*”

Computers like numbers: *75.102.246.202*

DNS: Domain Name Service

Given name get number (or vice versa)

Computers refer to themselves by the “loopback address”

127.0.0.1 (integers)

localhost or *localhost.localdomain* (string)

Ports:

Ports: *Which service on the given computer?*

Can range from 0 . . . 65535?

Common ones:

20 (ftp data), 21 (ftp control)

22 (ssh), 23 (telnet <-- DO NOT USE TO LOG IN!)

37 (time)

80 (www/http)

Packets (or datagrams)

- Any communication is split up into manageable chunks (“packets”) that are sent individually.
- These chunks get routing, checksum, cryptographic, *etc.* info added to them

Packet from layer N+1




Layer N *“Let’s compute the cryptographic hash and append it”*




Layer N-1 appends its own header




So I know with whom I wanna talk, how do I communicate? (1)



So I know with whom I wanna talk, how do I communicate? (2)



So I know with whom I wanna talk, how do I communicate? (3)



Tell me more about packets being sent over the 'net!

Two common networking protocols:

UDP: User Datagram Protocol

Fast! Unreliable!

“Just send the packets! Missed the last one? Don't worry, here's another!”

Question: When would you use this?

TCP: Transmission Control Protocol

Slower! Transmission is verified!

“Just a received a packet. Slow down! Let's make sure it hasn't been corrupted and that I have them all.”

Question: When would you use this?

Sockets!

Question: Cool deal! How do I program it?

Answer: With sockets, silly!

Socket communication

- Done with **read()** and **write()** (just like for files)
- Both have their own <address,port> pair
- Socket provides non-transient 2-way communication link

Server Side:

```
#include <sys/socket.h> //For socket ()
#include <netinet/in.h> //For sockaddr_in and htons ()
#include <netdb.h>       //For getaddrinfo ()
#include <errno.h>       //For errno var
#include <sys/stat.h>    //For open (), read (), write ()
#include <fcntl.h>       // and close ()
```

- **socket ()**: Ask OS for a socket
- **bind ()**: Bind socket and port together
- **listen ()**: Tell how many clients may queue
- **accept ()**: Wait until a client connects
- **write ()**: Write to client/server
- **read ()**: Read from client/server
- **close ()**: Close socket with client/server.

Client Side:

```
#include <sys/socket.h> //For socket ()
#include <netinet/in.h> //For sockaddr_in and htons ()
#include <netdb.h>       //For getaddrinfo ()
#include <errno.h>       //For errno var
#include <sys/stat.h>    //For open (), read (), write ()
#include <fcntl.h>       // and close ()
```

- **getaddrinfo ()**: Find server's IP address
- **socket ()**: Ask OS for a socket
- **connect ()**: Attempt to connect to server
- **write ()**: Write to server
- **read ()**: Read from server
- **close ()**: Close socket with server.

socket()

```
// Create a socket
int socketDescriptor =
    socket(AF_INET,          // AF_INET domain
          SOCK_STREAM,      // Reliable TCP
          0);
```

- Returns
 - A file descriptor that the server uses to see if a client has connected, or,
 - **-1** on error
- There's also **SOCK_DGRAM** for UDP
- Last parameter type if used for **SOCK_RAW**

bind()

// Bind socket to port

```
// We'll fill in this datastruct
struct sockaddr_in socketInfo;
```

```
// Fill socketInfo with 0's
```

```
memset(&socketInfo, '\0', sizeof(socketInfo));
```

```
// Use std TCP/IP
```

```
socketInfo.sin_family = AF_INET;
```

```
// Tell port in network endian with htons()
```

```
socketInfo.sin_port = htons(portNumber);
```

```
    // (1) Allow connections from myself only:
```

```
    struct in_addr addr;
```

```
    if (inet_aton("127.0.0.1", &addr) == 0)    exit(EXIT_FAILURE);
```

```
    socketInfo.sin_addr.s_addr = addr.s_addr;
```

```
    // or (2) Allow machine to connect to this service
```

```
    socketInfo.sin_addr.s_addr = INADDR_ANY;
```

```
// Try to bind socket with port and other specifications
```

```
int status = bind(socketDescriptor, // from socket()
```

```
                (struct sockaddr*)&socketInfo,
```

```
                sizeof(socketInfo));
```

```
status == -1 on error
```


What are those structs?

```
typedef uint32_t in_addr_t;
```

```
struct in_addr  
{  
    in_addr_t s_addr;  
};
```

```
struct sockaddr_in  
{  
    sa_family_t    sin_family; // addr family: AF_INET  
    in_port_t      sin_port;   // port (in network  
                                // byte order)  
    struct in_addr sin_addr;    // internet addr  
};
```

listen()

```
// Tell OS how many clients may queue  
// up for this server  
int status =  
    listen(socketDescriptor,  
           maxNumPendingClients);
```

- (Almost) ready to listen to port!
- 5 is a good default for `maxNumPendingClients`.
- If `status == -1` then error

accept()

```
// Accept connection to client  
int clientDescriptor =  
    accept(socketDescriptor, NULL, NULL) ;
```

- Wait (by default) for someone to actual connect
- Returns
 - a file descriptor for talking with one particular client, or
 - `-1` for error
- **connectionDescriptor** for talking with that one client (there may be others for other clients)
- **socketDescriptor** is for listening to socket.

Your turn!

Question: Hey! How is the server supposed to do two (or more!) things at once?

How do we get the server to both:

1. wait for another client to connect by listening to **socketDescriptor**, and
2. handle the current client(s) request by talking on **clientDescriptor**?

Do you speak *BIG* or *little* Endian?

Now that we're talking . . . we'd better use same endian!

```
// Host to network long (ie. 32-bit)
uint32_t htonl(uint32_t hostlong);

// Host to network short (ie. 16-bit)
uint16_t htons(uint16_t hostshort);

// Network to host long (ie. 32-bit)
uint32_t ntohl(uint32_t netlong);

// Network to host short (ie. 16-bit)
uint16_t ntohs(uint16_t netshort);
```

Your turn again!

- ◆ Write a server program that
 1. waits for a client to connect
 2. for any connected client it **read()**s characters and **write()**s the **toupper()** of them.

Client-side time!

getaddrinfo()

```
// Get info on server (the "hostName")
int getaddrinfo
(const char*    hostName, // e.g. "www.depaul.edu"
 const char*    service, // e.g. "ftp"
 const struct addrinfo* hints,
 struct addrinfo** resultPtr);
```

// Also: getnameinfo()
Gets info on host given integers

- Sets `resultPtr` to datastructure with info on host `hostName`
 - `hostName/service` effective tell `host:port`
 - returns integer: 0 == success, 0 != error.

Another strange struct!

```
struct addrinfo
{
    int             ai_flags;        // Used in hints
    int             ai_family;       // AF_INET, AF_INET6
                                     // or AF_UNSPEC for
                                     // IPv4, IPv6 or both
    int             ai_socktype;     // SOCK_STREAM
                                     // or SOCK_DGRAM
    int             ai_protocol;     // 0 = any protocol
    socklen_t       ai_addrlen;      // Len of next field
    struct sockaddr *ai_addr;        // (See prev slide)
    char            *ai_canonname;   // Official hostname
    struct addrinfo *ai_next;        // For linked list
};
```

getaddrinfo() (derived from wikipedia example)

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
#include <netdb.h>
#include <netinet/in.h>
#include <sys/socket.h>

#ifdef NI_MAXHOST
#define NI_MAXHOST 1025
#endif

#define LINE_LEN 256

#define SERVICE_ADDR_SEP_STR "://"

#define SERVICE_ADDR_SEP_STR_LEN \
    (sizeof(SERVICE_ADDR_SEP_STR)-1)
```

getaddrinfo() (derived from wikipedia example)

```
char*   enterUrlName(char* urlName,
                    int   urlNameLen
                    )
{
    printf("\n"
           "URL "
           "(e.g. ftp://ctilinux1.cstcis.cti.depaul.edu)"
           " or a blank line to stop\n"
           ": "
           );
    fgets(urlName,urlNameLen,stdin);

    char*   cPtr   = strchr(urlName,'\n');

    if (cPtr != NULL)
        *cPtr = '\0';

    return(urlName);
}
```

getaddrinfo() (derived from wikipedia example)

```
void parse (char* serviceName,
            int serviceNameLen,
            char* addrName,
            int addrNameLen,
            const char* urlName
            )
{
    const char* cPtr;
    for (cPtr = urlName; isspace(*cPtr); cPtr++);

    const char* sepPtr = strstr(cPtr, SERVICE_ADDR_SEP_STR);
    if (sepPtr == NULL)
    {
        strncpy(serviceName, "", serviceNameLen);
        strncpy(addrName, cPtr, addrNameLen);
    }
    else
    {
        int numServiceChars = sepPtr - cPtr;
        strncpy(serviceName, cPtr, numServiceChars);
        serviceName[numServiceChars] = '\0';
        strncpy(addrName, sepPtr + SERVICE_ADDR_SEP_STR_LEN, addrNameLen);
    }
}
```

getaddrinfo() (derived from wikipedia example)

```
void describe (const char* serviceName,
               const char* addrName
               )
{
    struct addrinfo* hostPtr;
    struct addrinfo* run;
    int status = getaddrinfo
        (addrName,
         (serviceName[0] == '\0')
         ? NULL
         : serviceName,
         NULL,
         &hostPtr
        );

    if (status != 0)
    {
        fprintf(stderr, gai_strerror(status));
        return;
    }
}
```

getaddrinfo() (derived from wikipedia example)

```
for (run = hostPtr; run != NULL; run = run->ai_next)
{
    char hostname[NI_MAXHOST] = "";

    int error = getnameinfo
                (run->ai_addr,
                 run->ai_addrlen,
                 hostname, NI_MAXHOST, NULL, 0, 0);

    if (error != 0)
    {
        fprintf(stderr, "error in getnameinfo: %s\n",
                gai_strerror(error));
        continue;
    }

    if (*hostname == '\0')
        printf("%-32s:", run->ai_canonname);
    else
        printf("%-32s:", hostname);
}
```

getaddrinfo() (derived from wikipedia example)

```
switch (run->ai_family)
{
case AF_INET :      printf("  (IPv4,"); break;
case AF_INET6 :     printf("  (IPv6,"); break;
case AF_UNSPEC :     printf("  (IPv4 & IPv6,"); break;
case AF_UNIX :      printf("  (local Unix,"); break;
case AF_IPX :        printf("  (Novell,"); break;
case AF_APPLETALK:   printf("  (Appletalk,"); break;
case AF_PACKET:      printf("  (Lo-level packet,"); break;
default :            printf("  (Unknown family?,");
}

switch (run->ai_socktype)
{
case SOCK_STREAM :printf("  TCP)\n"); break;
case SOCK_DGRAM : printf("  UDP)\n"); break;
case SOCK_SEQPACKET:printf("  sequenced, reliable)\n"); break;
case SOCK_RAW :    printf("  raw network protocol)\n"); break;
case SOCK_RDM :    printf("  reliable w/o ordering)\n"); break;
default :          printf("  unknown protocol?)\n");
}
}
}
```

getaddrinfo() (derived from wikipedia example)

```
int    main        ()
{
    char urlName[LINE_LEN];

    while ( *enterUrlName(urlName,LINE_LEN) != '\0' )
    {
        char  serviceName[LINE_LEN];
        char  addressName[LINE_LEN];

        parse(serviceName,LINE_LEN,addressName,LINE_LEN,urlName);
        describe(serviceName,addressName);
    }

    return(EXIT_SUCCESS);
}
```


connect()

```
// Connect to server
sockaddr_in server;

// Clear server datastruct
memset(&server, 0, sizeof(server));
// Use TCP/IP
server.sin_family      = AF_INET;
// Tell port # in proper network byte order
server.sin_port        = htons(portNumber);
// Copy connectivity info from info on server ("hostPtr")
server.sin_addr.s_addr =
    ((struct sockaddr_in*)hostPtr->ai_addr)->sin_addr.s_addr;

int status = connect(socketDescriptor, &server, sizeof(server));
```

➡ -1 means error

read(), write() and close()

As previously stated:

```
// Read from file/socket
// numRead==0 means "EndOfFile", numRead==-1 means "error"
int numRead =
    read(connectDescriptor,bufferAddress,bufferLen);
int numRead =
    recv(connectDescriptor,bufferAddress,bufferLen,int flags);

// Write to file/socket: numWritten == -1 means "error"
int numWritten =
    write(connectDescriptor,bufferAddress,bufferLen);
int numWritten =
    send(connectDescriptor,bufferAddress,bufferLen,int flags);

// Close connection: status == -1 means "error"
int status = close(descriptor);
```

But sometimes you don't want to wait for socket input

```
int recv(int connectDescriptor, void*  
    bufferPtr, int bufferLen, int flags)
```

Reads up to `bufferLen` bytes into the buffer pointed to by `bufferPtr` from file descriptor `connectDescriptor`. `flags` tells how to read, where `MSG_DONTWAIT` means "non-blocking".

Returns number of bytes read, or returns `-1` and sets global var `errno` to `EAGAIN` if the flag was `MSG_DONTWAIT` and there was nothing to read.

Short counts with `recv()`

- Short counts occur during:
 - Encounter End of file (EOF) when reading file (expected)
 - Reading text from terminal (also expected)
 - Reading from network or pipes if get interrupted by catching any sort of signal (an annoyance!)
 - **Question:** Did the fact that `read()` or `recv()` returned something mean that it **got** something, or that it ***was interrupted?***
- ***Oh no! Can nothing save us?!?!***

Robust I/O Package to the rescue!

```
/* From authors' thread-safe, buffered I/O package.
Same interface as read() */
ssize_t rio_readn(int fd, void *usrbuf, size_t n)
{
    size_t  nleft = n;
    ssize_t nread;
    char*    bufp = usrbuf;

    while (nleft > 0) {
        if ((nread = read(fd, bufp, nleft)) < 0) {
            if (errno==EINTR) /* Interrupted by sig handler rtn? */
                nread = 0;    /* Yes: Call read() again */
            else
                return -1;    /* No: Have some other error */
        }
        else if (nread == 0) /* Have EOF? */
            break;           /* Yes: Just quit loop */
        nleft -= nread;      /* Else that many fewer chars to get */
        bufp += nread;       /* Advance in buffer to read more */
    }
    return (n - nleft);     /* For non errors return val >= 0 */
}
```

Your turn!

- ◆ Write a client program that:
 1. Connects with the server
 2. Asks the user for text
 3. Sends the text to the server
 4. Gets the response back from the server and prints it

Next time

ncurses cursor control